

Emotion Detection Using Text

Import Libraries

```
[156]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import accuracy_score
from gensim.models import KeyedVectors
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Embedding, LSTM, Bidirectional, \
↳ Dropout
from tensorflow.keras.callbacks import EarlyStopping
from tensorflow.keras.preprocessing.sequence import pad_sequences
from sklearn.metrics import classification_report
from tensorflow.keras.models import save_model
from wordcloud import WordCloud
import pickle
import nltk
import warnings
from nltk.corpus import stopwords

nltk.download('stopwords')

warnings.filterwarnings("ignore")

%matplotlib inline
```

```
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
```

Merging the Dataset

```
[130]: def concatenate_files(file_list, output_file):
    with open(output_file, 'w') as outfile:
        for file_name in file_list:
            with open(file_name, 'r') as infile:
```

```

        for line in infile:
            outfile.write(line)

file_list = ['train.txt', 'test.txt', 'val.txt']
output_file = 'dataset.txt'

concatenate_files(file_list, output_file)

print("Files concatenated successfully as 'dataset.txt'")

```

Files concatenated successfully as 'dataset.txt'

Load the Dataset

```

[131]: df = pd.read_csv("dataset.txt", delimiter=';', header=None, names=['Sentences', 'Target'])
        sentences = df['Sentences'].values
        emotions = df['Target'].values

```

Dataset Exploration

```

[79]: print(df.head())

```

	Sentences	Target
0	i didnt feel humiliated	sadness
1	i can go from feeling so hopeless to so damned...	sadness
2	im grabbing a minute to post i feel greedy wrong	anger
3	i am ever feeling nostalgic about the fireplac...	love
4	i am feeling grouchy	anger

```

[80]: print(df.shape)

```

(20000, 2)

```

[81]: print(df.info())

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20000 entries, 0 to 19999
Data columns (total 2 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   Sentences  20000 non-null  object
 1   Target     20000 non-null  object
dtypes: object(2)
memory usage: 312.6+ KB
None

```

```

[82]: df.isnull().sum()

```

```
[82]: Sentences    0
      Target      0
      dtype: int64
```

```
[83]: print(df['Target'].unique())
```

```
['sadness' 'anger' 'love' 'surprise' 'fear' 'joy']
```

```
[84]: print((df.Target.value_counts() / df.shape[0] * 100).round(2))
```

```
Target
joy      33.80
sadness  28.98
anger    13.54
fear     11.86
love      8.20
surprise  3.60
Name: count, dtype: float64
```

Distribution of Target Variable


```
[85]: plt.figure(figsize=(8, 4))

ax = sns.countplot(x='Target', data=df, palette='Set2', order=df['Target'].
    ↪value_counts().index)

for p in ax.patches:
    ax.annotate(format(p.get_height(), '.0f'),
                (p.get_x() + p.get_width() / 2., p.get_height()),
                ha = 'center', va = 'center',
                xytext = (0, 5),
                textcoords = 'offset points')

plt.xlabel('Emotions', fontsize=12)
plt.ylabel('Count', fontsize=10)
plt.title('Distribution of Target variable', fontsize=14)

plt.tight_layout()
plt.show()
```



```
[86]: df["Sentences_length"] = [len(i) for i in df["Sentences"]]
```

```
[87]: df.head(5)
```

```
[87]:
```

	Sentences	Target	\
0	i didnt feel humiliated	sadness	
1	i can go from feeling so hopeless to so damned...	sadness	
2	im grabbing a minute to post i feel greedy wrong	anger	
3	i am ever feeling nostalgic about the fireplac...	love	
4	i am feeling grouchy	anger	

	Sentences_length
0	23
1	108
2	48
3	92
4	20

Max and Min Sentence length


```
[88]: print(df['Sentences_length'].max())
print(df['Sentences_length'].min())
```

```
300
```

```
7
```

```
[89]: sns.kdeplot(x=df["Sentences_length"], hue=df["Target"])
```

```
[89]: <Axes: xlabel='Sentences_length', ylabel='Density'>
```

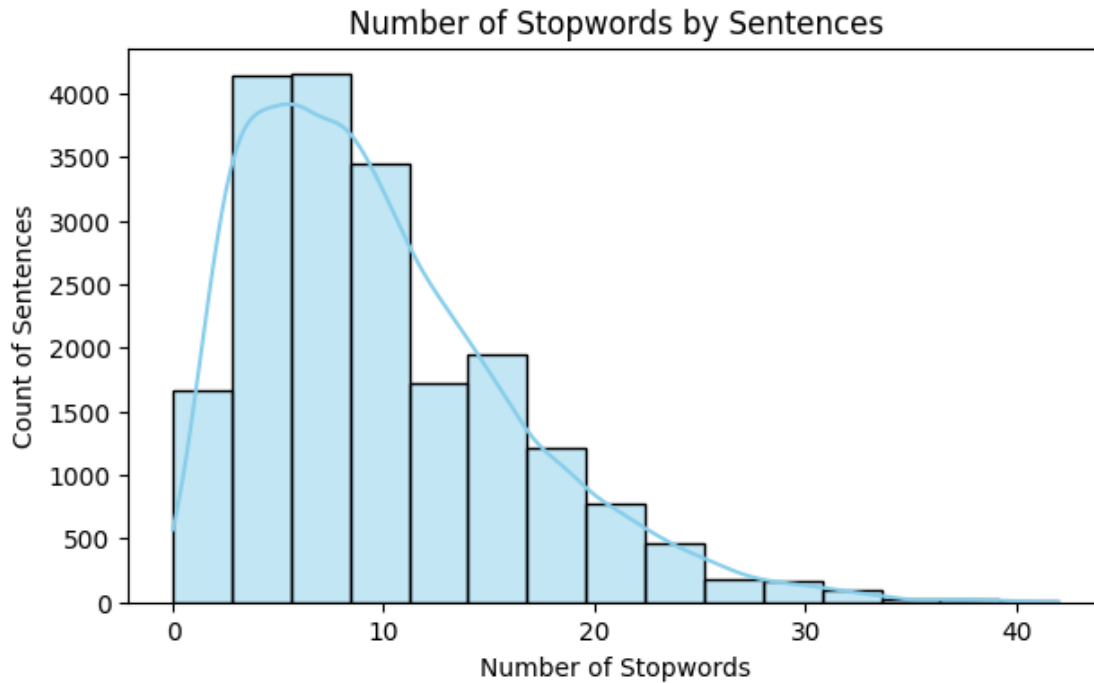


Number of Stopwords by Sentences

```
[105]: def count_stopwords(sentence):
        stop_words = set(stopwords.words('english'))
        tokens = sentence.split()
        return sum(1 for word in tokens if word.lower() in stop_words)

df['Stopword_Count'] = df['Sentences'].apply(count_stopwords)

plt.figure(figsize=(7, 4))
sns.histplot(df['Stopword_Count'], bins=15, kde=True, color='Skyblue')
plt.xlabel('Number of Stopwords')
plt.ylabel('Count of Sentences')
plt.title('Number of Stopwords by Sentences')
plt.show()
```



Removing the Stopwords

```
[106]: stop_words = set(stopwords.words('english'))
```


```
[107]: def remove_stopwords(text):
        tokens = text.split()
        filtered_tokens = [word for word in tokens if word.lower() not in
        ↪ stop_words]
        return ' '.join(filtered_tokens)

df['Sentences'] = df['Sentences'].apply(remove_stopwords)
```

Plot for words in Sentences after stopwords removal

```
[108]: sns.kdeplot(data=df, x=df["Sentences"].str.len(), hue=df["Target"])
```

```
[108]: <Axes: xlabel='Sentences', ylabel='Density'>
```



Word Cloud for Different Emotions

```
[109]: def generate_wordcloud(category_sentences, ax, title):
        wordcloud = WordCloud(width=400, height=300, background_color='white',
                                min_font_size=10).generate(' '.join(category_sentences))
        ax.imshow(wordcloud, interpolation='bilinear')
        ax.axis('off')
        ax.set_title(title)

category_groups = df.groupby('Target')['Sentences'].apply(list)

fig, axs = plt.subplots(2, 3, figsize=(12, 7))
fig.subplots_adjust(wspace=0.1)

for i, (category, sentences) in enumerate(category_groups.items()):
    row = i // 3
    col = i % 3
    generate_wordcloud(sentences, axs[row, col], category)

plt.suptitle('Word Cloud for Different Emotions', fontsize=16)
```

```
plt.show()
```

Word Cloud for Different Emotions



```
[116]: df.drop(["Sentences_length", "Stopword_Count"], axis = 1)
```

```
[116]:
```

		Sentences	Target
0		didnt feel humiliated	sadness
1	go feeling hopeless damned hopeful around some...		sadness
2	im grabbing minute post feel greedy wrong		anger
3	ever feeling nostalgic fireplace know still pr...		love
4		feeling grouchy	anger
...	
19995	im ssa examination tomorrow morning im quite w...		sadness
19996	constantly worry fight nature push limits inne...		joy
19997	feel important share info experience thing		joy
19998	truly feel passionate enough something stay tr...		joy
19999	feel like wanna buy cute make see online even one		joy

```
[20000 rows x 2 columns]
```

Downloading the Pretrained Model - GLoVe

```
[35]: !wget https://huggingface.co/stanfordnlp/glove/resolve/main/glove.6B.zip
```

--2024-05-07 14:45:48--


```

https://huggingface.co/stanfordnlp/glove/resolve/main/glove.6B.zip
Resolving huggingface.co (huggingface.co)... 65.8.178.12, 65.8.178.118,
65.8.178.93, ...
Connecting to huggingface.co (huggingface.co)|65.8.178.12|:443... connected.
HTTP request sent, awaiting response... 302 Found
Location: https://cdn-lfs.huggingface.co/stanfordnlp/glove/6471382cdd837544bf3ac
72497a38715e845897d265b2b424b4761832009c837?response-content-disposition=attachm
ent%3B+filename%3DUTF-8%27%27glove.6B.zip%3B+filename%3D%22glove.6B.zip%22%3B&r
esponse-content-type=application%2Fzip&Expires=1715352348&Policy=eyJTdGF0ZWlbnQ
iOlt7IkNvbmlRdGlzbiI6eyJEYXRlTGZvc1RoYW4iOnsiQVdT0kVWb2NoVGltZSI6MTcxNTM1MjM0OH1
9LCJSZXNvdXJzSI6Imh0dHBzOi8vY2RuLWxmcy5odWdnaW5nZmFjZS5jb3R5ZG5scC9nbG9
2ZS82NDcxMzgyY2RkODM3NTQ0YmZyZWY3MjQ5NDU4OTdkMjY1YjJiNDI0YjQ3NjE4MzI
wMDljODM3P3Jlc3BvbnNlLWVbnRlbnQtZGlzcG9zaXRpb249KiZyZXNwb25zZS1jb250ZW50LXR5cGU
9KiJ9XX0_&Signature=elvMx%7EY0z8LyztBDd3Ic-wEEGYpU1cZ12fP%7EPHj1BjzwlgrXgdmzqFve
fml%7EXAjR8%7E44xhBOWEoukJ9j1kn8smSkD4XMK1de-
CkkBQ%7EeddF1UgvG931SFSvNgntbN7AMxG9yDLvJZPWKh7xBQ2oIpr26LIKdrAhrjpHf-YmRwxKWZ7
HvMrU%7ECDP8pPyPRZVyB0xiPLavBTTVLFFxWyaPP43RbKQCw%7EYNqcwudMrbXdrjaXgJSTj3r4lwl
PZI%7EIOE7BI%7EzzoVGfSMs0kS0kcYcVZmoLkXWoRuqhHGWBlykGRPCVKtUMU12VeuhAwX%7Efpjd5X
Xlwb2REKvPVo8IV7Q_&Key-Pair-Id=KVTP0A1DKRTAX [following]
--2024-05-07 14:45:48-- https://cdn-lfs.huggingface.co/stanfordnlp/glove/647138
2cdd837544bf3ac72497a38715e845897d265b2b424b4761832009c837?response-content-disp
osition=attachment%3B+filename%3DUTF-8%27%27glove.6B.zip%3B+filename%3D%22glove
.6B.zip%22%3B&response-content-type=application%2Fzip&Expires=1715352348&Policy=
eyJTdGF0ZWlbnQlOlt7IkNvbmlRdGlzbiI6eyJEYXRlTGZvc1RoYW4iOnsiQVdT0kVWb2NoVGltZSI6
MTcxNTM1MjM0OH19LCJSZXNvdXJzSI6Imh0dHBzOi8vY2RuLWxmcy5odWdnaW5nZmFjZS5jb3R5ZG5scC9nbG92ZS82NDcxMzgyY2RkODM3NTQ0YmZyZWY3MjQ5NDU4OTdkMjY1YjJi
NDI0YjQ3NjE4MzIwMDljODM3P3Jlc3BvbnNlLWVbnRlbnQtZGlzcG9zaXRpb249KiZyZXNwb25zZS1j
b250ZW50LXR5cGU9KiJ9XX0_&Signature=elvMx%7EY0z8LyztBDd3Ic-wEEGYpU1cZ12fP%7EPHj1B
jzwlgrXgdmzqFvefml%7EXAjR8%7E44xhBOWEoukJ9j1kn8smSkD4XMK1de-
CkkBQ%7EeddF1UgvG931SFSvNgntbN7AMxG9yDLvJZPWKh7xBQ2oIpr26LIKdrAhrjpHf-YmRwxKWZ7
HvMrU%7ECDP8pPyPRZVyB0xiPLavBTTVLFFxWyaPP43RbKQCw%7EYNqcwudMrbXdrjaXgJSTj3r4lwl
PZI%7EIOE7BI%7EzzoVGfSMs0kS0kcYcVZmoLkXWoRuqhHGWBlykGRPCVKtUMU12VeuhAwX%7Efpjd5X
Xlwb2REKvPVo8IV7Q_&Key-Pair-Id=KVTP0A1DKRTAX
Resolving cdn-lfs.huggingface.co (cdn-lfs.huggingface.co)... 108.157.173.21,
108.157.173.84, 108.157.173.44, ...
Connecting to cdn-lfs.huggingface.co (cdn-
lfs.huggingface.co)|108.157.173.21|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 862182753 (822M) [application/zip]
Saving to: 'glove.6B.zip'

```

```

glove.6B.zip          100%[=====>] 822.24M  58.8MB/s    in 12s

```

```

2024-05-07 14:46:00 (69.1 MB/s) - 'glove.6B.zip' saved [862182753/862182753]

```

```
[36]: import zipfile
zip_ref = zipfile.ZipFile("glove.6B.zip", 'r')
zip_ref.extractall(".")
zip_ref.close()
```

Load the model

```
[37]: def load_glove_model(File):
    print("Loading Glove Model")
    glove_model = {}
    with open('glove.6B.300d.txt', 'r') as f:
        for line in f:
            split_line = line.split()
            word = split_line[0]
            embedding = np.array(split_line[1:], dtype=np.float64)
            glove_model[word] = embedding
    print(f"{len(glove_model)} words loaded!")
    return glove_model
```


```
[38]: glove_model = load_glove_model('glove.6B.300d.txt')
```

Loading Glove Model
400001 words loaded!

Tokenize the sentences and create word-to-index mapping

```
[161]: word_to_index = {}
index = 1 # Start index from 1, leaving 0 for padding
for sentence in sentences:
    for word in sentence.split():
        if word not in word_to_index:
            word_to_index[word] = index
            index += 1
```

Creating an embedding matrix

```
[136]: embedding_matrix = np.zeros((len(word_to_index) + 1, len(glove_model['the']))) 
    ↪ # Add 1 to include padding token
    for word, index in word_to_index.items():
        if word in glove_model:
            embedding_matrix[index] = glove_model[word]
```

Performing label encoding on the target labels

```
[137]: label_encoder = LabelEncoder()
encoded_emotions = label_encoder.fit_transform(emotions)
```

Split the dataset into training, validation, and testing sets

```
[138]: X_train, X_test, y_train, y_test = train_test_split(sentences,
↳ encoded_emotions, test_size=0.2, random_state=42)
X_train, X_val, y_train, y_val = train_test_split(X_train, y_train, test_size=0.
↳ 1, random_state=42)
```

Converting sentences to sequences of indices

```
[139]: X_train_sequences = [[word_to_index[word] for word in sentence.split() if word
↳ in word_to_index] for sentence in X_train]
X_val_sequences = [[word_to_index[word] for word in sentence.split() if word in
↳ word_to_index] for sentence in X_val]
X_test_sequences = [[word_to_index[word] for word in sentence.split() if word
↳ in word_to_index] for sentence in X_test]
```

Padding sequences to make them of equal length

```
[140]: max_sequence_length = max(len(sequence) for sequence in X_train_sequences +
↳ X_val_sequences + X_test_sequences)
X_train_padded = pad_sequences(X_train_sequences, maxlen=max_sequence_length,
↳ padding='post')
X_val_padded = pad_sequences(X_val_sequences, maxlen=max_sequence_length,
↳ padding='post')
X_test_padded = pad_sequences(X_test_sequences, maxlen=max_sequence_length,
↳ padding='post')
```

Train the Model

```
[141]: bilstm_model = Sequential([
    Embedding(input_dim=len(word_to_index) + 1,
↳ output_dim=len(glove_model['the']), weights=[embedding_matrix],
↳ input_length=max_sequence_length, trainable=False),
    Bidirectional(LSTM(64, dropout=0.2, recurrent_dropout=0.2)),
    Dropout(0.2), # Dropout layer added after the Bidirectional LSTM layer
    Dense(len(label_encoder.classes_), activation='softmax')
])
bilstm_model.compile(loss='sparse_categorical_crossentropy', optimizer='adam',
↳ metrics=['accuracy'])

# Early stopping criteria
early_stopping = EarlyStopping(monitor='val_loss', patience=3,
↳ restore_best_weights=True)
```

```
[159]: bilstm_model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		

embedding (Embedding)	(None, 66, 300)	5129100
bidirectional (Bidirectional)	(None, 128)	186880
dropout (Dropout)	(None, 128)	0
dense (Dense)	(None, 6)	774

```

=====
Total params: 5316754 (20.28 MB)
Trainable params: 187654 (733.02 KB)
Non-trainable params: 5129100 (19.57 MB)
-----

```

```

[142]: bilstm_history = bilstm_model.fit(X_train_padded, y_train, epochs=25,
    ↪batch_size=32, validation_data=(X_val_padded, y_val),
    ↪callbacks=[early_stopping])

```

```

Epoch 1/25
450/450 [=====] - 172s 369ms/step - loss: 1.2147 -
accuracy: 0.5519 - val_loss: 0.7651 - val_accuracy: 0.7163
Epoch 2/25
450/450 [=====] - 166s 368ms/step - loss: 0.6749 -
accuracy: 0.7642 - val_loss: 0.4767 - val_accuracy: 0.8300
Epoch 3/25
450/450 [=====] - 219s 488ms/step - loss: 0.4701 -
accuracy: 0.8383 - val_loss: 0.3492 - val_accuracy: 0.8662
Epoch 4/25
450/450 [=====] - 149s 332ms/step - loss: 0.3659 -
accuracy: 0.8722 - val_loss: 0.3035 - val_accuracy: 0.8869
Epoch 5/25
450/450 [=====] - 160s 357ms/step - loss: 0.3098 -
accuracy: 0.8865 - val_loss: 0.2732 - val_accuracy: 0.8950
Epoch 6/25
450/450 [=====] - 146s 325ms/step - loss: 0.2547 -
accuracy: 0.9043 - val_loss: 0.2493 - val_accuracy: 0.9050
Epoch 7/25
450/450 [=====] - 153s 341ms/step - loss: 0.2297 -
accuracy: 0.9133 - val_loss: 0.2251 - val_accuracy: 0.9069
Epoch 8/25
450/450 [=====] - 141s 313ms/step - loss: 0.2001 -
accuracy: 0.9233 - val_loss: 0.2120 - val_accuracy: 0.9100
Epoch 9/25
450/450 [=====] - 152s 337ms/step - loss: 0.1865 -
accuracy: 0.9272 - val_loss: 0.2197 - val_accuracy: 0.9112
Epoch 10/25
450/450 [=====] - 146s 325ms/step - loss: 0.1678 -

```

```

accuracy: 0.9331 - val_loss: 0.2095 - val_accuracy: 0.9150
Epoch 11/25
450/450 [=====] - 157s 348ms/step - loss: 0.1530 -
accuracy: 0.9408 - val_loss: 0.2150 - val_accuracy: 0.9081
Epoch 12/25
450/450 [=====] - 151s 336ms/step - loss: 0.1455 -
accuracy: 0.9410 - val_loss: 0.2093 - val_accuracy: 0.9119
Epoch 13/25
450/450 [=====] - 144s 321ms/step - loss: 0.1328 -
accuracy: 0.9447 - val_loss: 0.2094 - val_accuracy: 0.9131
Epoch 14/25
450/450 [=====] - 151s 335ms/step - loss: 0.1226 -
accuracy: 0.9520 - val_loss: 0.2132 - val_accuracy: 0.9169
Epoch 15/25
450/450 [=====] - 150s 333ms/step - loss: 0.1134 -
accuracy: 0.9532 - val_loss: 0.2303 - val_accuracy: 0.9169

```

Evaluating the model

```

[143]: loss, accuracy = bilstm_model.evaluate(X_test_padded, y_test)
print("BiLSTM Accuracy:", accuracy)

```

```

125/125 [=====] - 5s 41ms/step - loss: 0.2327 -
accuracy: 0.9062
BiLSTM Accuracy: 0.90625

```

```

[145]: y_pred_prob = bilstm_model.predict(X_test_padded)
y_pred = np.argmax(y_pred_prob, axis=1)

print("Classification Report:")
print(classification_report(y_test, y_pred, target_names=label_encoder.
↪classes_))

```

```

125/125 [=====] - 18s 122ms/step


```

Classification Report:

	precision	recall	f1-score	support
anger	0.90	0.93	0.91	536
fear	0.86	0.83	0.84	458
joy	0.94	0.92	0.93	1339
love	0.78	0.84	0.81	335
sadness	0.95	0.94	0.95	1173
surprise	0.73	0.85	0.78	159
accuracy			0.91	4000
macro avg	0.86	0.88	0.87	4000
weighted avg	0.91	0.91	0.91	4000


Plot for training and validation loss

```
[146]: plt.plot(bilstm_history.history['loss'], label='Training Loss')
plt.plot(bilstm_history.history['val_loss'], label='Validation Loss')
plt.title('Training and Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```



Plot for training and validation accuracy

```
[147]: plt.plot(bilstm_history.history['accuracy'], label='Training Accuracy')
plt.plot(bilstm_history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Training and Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.show()
```



Predicting New Sentences

```
[154]: new_sentences = [
    "I feel excited about the upcoming project",
    "This situation makes me anxious",
    "I am calm and relaxed right now"
]

tokenizer = lambda x: [[word_to_index[word] for word in sentence.split() if
    ↪word in word_to_index] for sentence in x]

new_sequences = tokenizer(new_sentences)
new_sequences_padded = pad_sequences(new_sequences, maxlen=max_sequence_length,
    ↪padding='post')

predictions = bilstm_model.predict(new_sequences_padded)

predicted_emotions = label_encoder.inverse_transform(np.argmax(predictions,
    ↪axis=1))
```

```
print("Predicted Emotions for New Sentences:")
for sentence, emotion in zip(new_sentences, predicted_emotions):
    print(f"Sentence: {sentence} | Predicted Emotion: {emotion}")
```

1/1 [=====] - 1s 1s/step

Predicted Emotions for New Sentences:

Sentence: I feel excited about the upcoming project | Predicted Emotion: joy

Sentence: This situation makes me anxious | Predicted Emotion: fear

Sentence: I am calm and relaxed right now | Predicted Emotion: joy

Save Model

```
[157]: bilstm_model.save('bilstm_model.h5')

with open('bilstm_model.pkl', 'wb') as f:
    pickle.dump(bilstm_model, f)

print("Model saved successfully!")
```

Model saved successfully!